

# Algolab BGL Introduction

Andreas Bärtschi

October 7, 2015



A **generic** C++ library of graph data structures and algorithms.

**BGL docs** – your new best friend:

[http://www.boost.org/doc/libs/1\\_57\\_0/libs/graph/doc](http://www.boost.org/doc/libs/1_57_0/libs/graph/doc)

Moodle: There's a brief **copy & paste manual**.

# BGL: A generic library

Genericity type	STL	BGL
Algorithm / Data-Structure Interoperability	Decoupling of algorithms and data-structures Key ingredients: iterators	Decoupling of graph algorithms and graph representations Vertex iterators, edge iterators, adjacency iterators
Parameterization	Element type parameterization	Vertex and edge property multi-parametrization <i>Associate multiple properties</i> <i>Accessible via property maps</i>
Extensions	through function objects	through a <i>visitor object</i> , event points and methods depend on particular algorithm

# BGL: A generic library

Genericity type	STL	BGL
Algorithm / Data-Structure Interoperability	Decoupling of algorithms and data-structures Key ingredients: iterators	Decoupling of graph algorithms and graph representations Vertex iterators, edge iterators, adjacency iterators
Parameterization	Element type parameterization	Vertex and edge property multi-parametrization <i>Associate multiple properties</i> <i>Accessible via property maps</i>
Extensions	through function objects	through a <i>visitor object</i> , event points and methods depend on particular algorithm

# BGL: Graph Representations / Data Structures

Structure	Representation	Advantages	Do
Graph classes	Adjacency list	Swiss army knife: Directed/undirected graphs, allow/disallow parallel-edges, efficient insertion, fast adjacency structure exploitation	<b>use this!</b>
	Adjacency matrix	Dense graphs	<i>use at your</i>
Adaptors	Edge list	Simplicity	<i>own risk!</i>
	External adaptation	Convert existing graph structures (LEDA etc.) to BGL	Not covered in Algotlab.

## BGL: adjacency\_list

Example **without** any vertex or edge properties:

```
// Easy syntax
```

```
typedef adjacency_list<vecS, vecS, directedS>    Graph;
```

```
...
```

```
// which is the same as:
```

```
typedef adjacency_list<vecS, vecS, directedS,  
                      no_property,  
                      no_property>    Graph;
```

```
...
```

## BGL: adjacency\_list

Example **with** vertex property and multiple edge properties:

```
// Note syntax for defining more than one edge property
typedef adjacency_list<vecS, vecS, directedS,
    property<vertex_name_t, string>, // vertex property
    property<edge_capacity_t, int,    // multiple edge properties: nested
        property<edge_residual_capacity_t, int,
            property<edge_reverse_t, Traits::edge_descriptor> > > > Graph;

typedef property_map<Graph, vertex_name_t>::type      NameMap;
typedef property_map<Graph, edge_capacity_t>::type    CapacityMap;
typedef property_map<Graph, edge_residual_capacity_t>::type ResidualMap;
typedef property_map<Graph, edge_reverse_t>::type     ReverseMap;
...
```

# BGL: Graph Algorithms

Area	Topic	Details
Basics	Distances	Dijkstra shortest paths Prim minimum spanning tree Kruskal minimum spanning tree
	Components	Connected, biconnected & strongly connected components
	General Matchings	General unweighted matching
Flows	Maximum Flow	Graph setup (residual graph) Edmonds-Karp and Push-Relabel
	Disjoint paths	Vertex- / Edge-disjoint s-t paths
Advanced Flows	Minimum Cut	Maxflow-Mincut Theorem
	Bipartite Matchings	Vertex Cover & Independent Set
	Mincost Maxflow	Bipartite weighted matching & more

Many more (not in Algotab 2015): planarity testing, sparse matrix ordering, ...

**Prerequisites:** Theory, BFS, DFS, topological sorting, Eulerian tours, Union-Find...



## Tutorial problem: statement


**Input:** A directed graph  $G$  with positive weights on edges and a vertex  $t$ .

**Definition:** We call a vertex  $u$  *universal* if all vertices in  $G$  can be reached from it.


**Output:** The length of a shortest path  $u \rightarrow t$  that starts in some universal vertex  $u$ . If such a path does not exist, output NO.

$$|V(G)| \leq 10^5, |E(G)| \leq 2 \cdot 10^5.$$

## Tutorial problem: example



## Tutorial problem: example



## Tutorial problem: how to start?

Time's short, so hurry up!


- "Check if there is a unique  $u$  with no in-edges, if yes output shortest path  $u \rightarrow t$ ."  
(**what if there is no such  $u$ ?**)
- "For each  $u$  check with DFS if  $u$  reaches all vertices, then..." (**too slow**)
- `#include <iostream>`  
`int main()`  
`// some random algorithm`

**No! Take your time,**  
model the problem,  
design the algorithm,  
**understand why it should work,**  
 $\Rightarrow$  then code.


## Tutorial problem: how to start?

- Bad question: *Why shouldn't it work?*  
("It is correct on all three examples I came up with", etc.)
- Good question: *Why should it work?*  
("How would I prove it works?")
- Applies to Moodle forums as well!

## Tutorial problem: example




## Tutorial problem: example




⇒ must be related to some sort of connected component concept in directed graphs!

## Tutorial problem: strongly connected components (SCC) example






## Tutorial problem: SCC directed acyclic graph example



Is there always a universal vertex?

## Tutorial problem: strongly connected components “NO” example

No!



## Tutorial problem: modeling

Let us call a strongly connected component with no in-edges in the SCC Directed Acyclic Graph a *minimal component*.

### Fact

*If there is more than one minimal component in  $G$ , then there is no universal  $u$ .*

### Lemma

*If there is exactly one minimal component in  $G$ , then its vertices are exactly the universal vertices.*

## Tutorial problem: modeling

New formulation of the problem:

- 1 If there exists  $> 1$  minimal strongly connected component in  $G$ , output NO.
- 2 Output the shortest distance  $u \rightarrow t$  for universal  $u$  in  $G$ .

But this is still  $\Omega(n^2)$  in the worst case.

## Tutorial problem: modeling

New formulation of the problem:

- 1 If there exists  $> 1$  maximal strongly connected component in  $G_T$ , output NO.
- 2 Output the shortest distance  $t \rightarrow u$  for universal  $u$  in  $G_T$ .

Now we can work only with  $G_T$ .

# Tutorial problem: implementation

First and foremost, [BGL docs](#):

- How to find the [strong\\_components](#).
- How to check how many maximal components are there?  
[topological\\_sort](#)?  
Maybe there is a simple ad hoc?
- Compute shortest  $t - u$  path on  $G_T$  with [dijkstra\\_shortest\\_paths](#).

## Tutorial problem: code – preamble

```
1: #include <climits>
2: #include <iostream>
3: #include <vector>
4:
5: #include <boost/graph/adjacency_list.hpp>
6: #include <boost/graph/dijkstra_shortest_paths.hpp>
7: #include <boost/graph/strong_components.hpp>
8: #include <boost/tuple/tuple.hpp> // tuples::ignore
9:
10: using namespace std;
11: using namespace boost;
```

## Tutorial problem: code – typedefs

```
13: // Directed graph with int weights on edges.
14: typedef adjacency_list<vecS, vecS, directedS,
15:         no_property,
16:         property<edge_weight_t, int> > Graph;
17: typedef graph_traits<Graph>::edge_descriptor   Edge;    // Edge type
18: typedef graph_traits<Graph>::edge_iterator      EdgeIt;   // Iterator
19: // Map edge -> weight.
20: typedef property_map<Graph, edge_weight_t>::type   WeightMap;
```



## Tutorial problem: code – main

```
22: void testcase();
23:
24: int main() {
25:     ios_base::sync_with_stdio(false);
26:     int T; // First input line: Number of testcases.
27:     cin >> T;
28:     while (T--) testcase();
29: }
```

## Tutorial problem: code – reading the input

```
31: void testcase() {
32:     int V, E, t; // 1st line: <ver_no> <edg_no> <tgt>
33:     cin >> V >> E >> t;
34:     Graph G(V);
35:     WeightMap wm = get(edge_weight, G);
36:     for (int i = 0; i < E; ++i) {
37:         int u, v, c;
38:         Edge e; // Each edge: <src> <tgt> <cost>
39:         cin >> u >> v >> c; // Swap u, v and instead of G
40:         tie(e, tuples::ignore) = add_edge(v, u, G); // create G_T!
41:         wm[e] = c;
42:     }
```

## Tutorial problem: code – strong components

```
31: void testcase() {  
    ...  
44:     // Store index of the vertices' strong component  
45:     vector<int> scc(V);  
46:     int nsc = strong_components(G,  
47:                               make_iterator_property_map(scc.begin(),  
48:                               get(vertex_index, G)));
```

## Tutorial problem: code – maximal SCCs

```
25: void testcase() {  
    ...  
50:    // Find universal strong component (if any)  
51:    // Why does this approach work? Exercise.  
52:    vector<int> is_max(nsc, 1);  
53:    EdgeIterator ebegin, end;  
54:    // Iterate over all edges.  
55:    for (tie(ebegin, end) = edges(G); ebegin != end; ++ebegin) {  
56:        // ebegin is an iterator, *ebegin is a descriptor.  
57:        int u = source(*ebegin, G), v = target(*ebegin, G);  
58:        if (scc[u] != scc[v]) is_max[scc[u]] = 0;  
59:    }
```

## Tutorial problem: code – maximal SCCs

```
25: void testcase() {  
    ...  
60:     int max_count = count(is_max.begin(), is_max.end(), true);  
61:     if (max_count != 1) {  
62:         cout << "NO" << endl;  
63:         return;  
64:     }
```


## Tutorial problem: code – Dijkstra

```
25: void testcase() {  
    ...  
66:    //Compute shortest t-u path in G_T  
67:    vector<int> dist(V);  
68:    vector<int> pred(V);  
69:    dijkstra_shortest_paths(G, t,  
70:        predecessor_map(make_iterator_property_map(pred.begin(),  
71:            get(vertex_index, G)))).  
72:        distance_map(make_iterator_property_map(dist.begin(),  
73:            get(vertex_index, G))));
```


## Tutorial problem: code – output

```
25: void testcase() {  
    ...  
74:     int res = INT_MAX;  
75:     for (int u = 0; u < V; ++u)  
76:         // Minimum of distances to 'maximal' vertices  
77:         if (is_max[scc[u]])  
78:             res = min(res, dist[u]);  
79:     cout << res << endl;  
80: }
```

# Minimum spanning trees



Prim Minimum Spanning Tree



Kruskal Minimum Spanning Tree



# Minimum spanning tree algorithms

Algorithm	starts with	next	Time
Prim MST	Arbitrary start vertex	Adds connection (if possible) to the closest neighbour of all so far discovered vertices.	$O(E \log V)$
Kruskal	Edge of minimum weight	Adds next smallest edge, if this is possible without creating a cycle.	$O(E \log E)$

We need to provide a predecessor vector to Prim (stores parents in the MST), and an edge vector to Kruskal (stores the edges of the MST).

# Minimum spanning tree implementations


## Prim's algorithm

```
vector<int> p(V); // predecessor vector
prim_minimum_spanning_tree(G, &p[0]);
for (int j = 0; j < V; ++j) {
    Edge e; bool success;
    tie(e, success) = edge(j, p[j], G);
    if (success) { // careful: doesn't work like this when G has loops
        ...
    }
}
```

## Kruskal's algorithm

```
vector<Edge> mst; // edge vector to store mst
kruskal_minimum_spanning_tree(G, back_inserter(mst));
vector<Edge>::iterator ebegin, eend = mst.end();
for (ebegin = mst.begin(); ebegin != eend; ++ebegin) {
    ...
}
```

## Maximum matching: general unweighted version



- $G = (V, E)$
- $M \subseteq E$  is a matching if and only if no two edges of  $M$  are adjacent.
- In an unweighted graph, a maximum matching is a matching of maximum cardinality.
- In a weighted graph, a maximum matching is a matching such that the weight sum over the included edges is maximum.
- BGL does not provide weighted matching algorithms.

## Maximum matching: invoking algorithm

```
#include <boost/graph/max_cardinality_matching.hpp>
:

vector<VertexDescriptor> mate(n);
edmonds_maximum_cardinality_matching(g, &mate[0]);

const VertexDescriptor NULL_VERTEX = graph_traits<Graph>::null_vertex();
...
for(int i = 0; i < n; i++)
    if(mate[i] != NULL_VERTEX && i < mate[i])
        cout << i << " - " << mate[i] << endl;
```

## Getting started: BGL installation

- Pre-installed in ETH computer rooms and the Algalab Virtualbox Image.  
Most likely also already installed on your system if you installed CGAL last week.
- On "standard" Linux distributions try getting a package from the repository.  
On MacOS package from [MacPorts](#).
- Comments on the versions:
  - 1.57: This version runs on the judge.
  - 1.55+: These versions have Mincost-maxflow, should be fine.
  - 1.54: Prim MST bug (unless Ubuntu)

## Getting started: BGL without installing

- BGL is a Header-only library.
- Download recent version from: <http://www.boost.org/users/download/>.
- Just unpack the .tar.bz2 file, no installation required, see Section 3 here: [http://www.boost.org/doc/libs/1\\_57\\_0/more/getting\\_started/unix-variants.html](http://www.boost.org/doc/libs/1_57_0/more/getting_started/unix-variants.html).
- To build using this version of boost use this command:  
*g++ -I path/to/boost\_1\_57\_0 example.cpp -o example*
- Explanation: The '-I' flag tells the compiler to include all the files from this directory, so that it can find header files like 'boost/graph/adjacency\_list.hpp'

## Getting started: compilation problems

Error messages can be terrible.

- Consider re-compiling the code after every line after it is first written. This will help to identify the problem quickly.
- Especially after the typedefs, and again after building the graph, before you do anything else!
- There will be confusing typedefs, nested types, iterators etc. Come up with a naming pattern and stick to it.

## Getting started: runtime problems

- Isolate the smallest possible example where the program misbehaves.
- Watch out for invalidated iterators.
- Print a graph and see if it looks as expected. In particular, check if the number of vertices didn't increase due to mistakes in your edge insertion.
- More on the slides of the next (and last) Section of today.



## Getting started: Using the forums

Some post	Good post
<p>I tried to solve this question as mentioned in the lecture slides, I got timelimit, I did not yet apply the</p> <p>Spoiler» sort to make it fast «, but in the slides it is mentioned that without</p> <p>Spoiler» sort «, it is still fast enough, I will be grateful if you could mention the problem with my code that makes it slow, thanks</p>	<p>My code to Problem xy gets a timelimit on the last test set and I don't know why. My approach was the following:</p> <p>Spoiler »</p> <p>«</p> <p>I can argue that my solution is correct, because</p> <p>Spoiler »</p> <p>«</p> <p>The overall running time of my solution is</p> <p>Spoiler»</p> <p>«</p> <p>Attached you can find my (reasonably commented) submission.</p>

## Getting started: Problem of the week


As usual, on Monday. Don't miss it!

Be advised it doesn't have to be BGL.

Anything already covered in the course can be used.



**BGL**  
THE BOOST GRAPH LIBRARY



## Useful stuff: Namespace

Simple: everything is in namespace boost.

```
using namespace boost;
```

## Useful stuff: Tuple utility functions

A handy trick when function returns a tuple:

```
#include <boost/tuple/tuple.hpp>
...
int a;
double b;
tie(a, b) = make_pair(1, 0.1);
tie(tuples::ignore, b) = make_pair(2, -0.3);
```

## Useful stuff: adjacency\_list

This is the class you almost always need.

```
#include <boost/graph/adjacency_list.hpp>
```

```
...
```

```
typedef adjacency_list<vecS, vecS, directedS> Graph;
```

- 1st vecS — for each vertex, adjacency list kept in a vector.  
Choosing setS instead disallows parallel edges.
- 2nd vecS — a list of all edges is kept in a vector.
- directedS — directed graph.  
Other choice: undirectedS.

## Useful stuff: Vertices and edges

- **Vertex descriptor:** This is (almost) always an int in range  $[0, \text{num\_vertices}(G))$ . Don't overcomplicate this.
- **Edge descriptor:** an object that represents a single edge.

```
typedef graph_traits<Graph>::edge_descriptor Edge;  
Edge e;  
int u = source(e, G), v = target(e, G);
```

## Useful stuff: Building a graph

```
Graph G(n);    // Constructs empty graph with n vertices
...
Edge e;
bool ok;
tie(e, ok) = add_edge(u, v, G);
```

- Adds edge from  $u$  to  $v$  in  $G$ .
- Caveat: if  $u$  or  $v$  don't exist in the graph,  $G$  is automatically extended.
- Returns an  $(\text{Edge}, \text{bool})$  pair. First coordinate is an edge descriptor. If parallel edges are allowed, second coordinate is always true. Otherwise it is false in case of failure when the edge is a duplicate.



## Useful stuff: Removing vertices and edges (dangerous!)

```
remove_edge(u, v, G);  
remove_edge(e, G);  
clear_vertex(u, G);  
clear_out_edges(u, G);  
remove_vertex(u, G);
```

- Consult the docs. Takes time, invalidates descriptors and iterators, might behave counterintuitively. Not recommended.

## Useful stuff: Clearing a graph

```
G.clear();
```

- Removes all edges and vertices.

```
G = Graph(n);
```

- Destroys the old graph and creates a new one with `n` vertices.

## Useful stuff: Iterating over vertices and edges

```
// Iterating over vertices
for (int u = 0; u < num_vertices(G); ++u) {
    ...
}

// Iterating over edges
typedef graph_traits<Graph>::edge_iterator EdgeIterator;
EdgeIterator eit, eend;
for (tie(eit, eend) = edges(G); eit != eend; ++eit) {
    // eit is EdgeIterator, *eit is EdgeDescriptor
    int u = source(*eit, G), v = target(*eit, G);
    ...
}
```

- `edges(G)` returns a pair of iterators which define a range of all edges.
- For undirected graphs each edge is visited once, with some orientation.

## Useful stuff: Iterating over edges incident to a vertex

```
// Iterating over outgoing edges
typedef graph_traits<Graph>::out_edge_iterator
    OutEdgeIterator;
OutEdgeIterator eit, eend;
for (tie(eit, eend) = out_edges(u, G); eit != eend; ++eit) {
    int v = target(*eit, G);
    ...
}
```

- `source(*eit, G)` is guaranteed to be `u`, even in an undirected graph.

## Useful stuff: Property maps

- Think of the *property map* as a map (i.e., object with operator `[]`) indexed by vertices or edges.
- Property maps of vertices can be simulated with a vector, but maps of edges are very convenient.

## Useful stuff: Vertex property map

```
// Note syntax for defining more than one map.
typedef adjacency_list<vecS, vecS, directedS,
    property<vertex_name_t, string,
        property<vertex_distance_t, int> > > Graph;
typedef property_map<Graph, vertex_name_t>::type NameMap;
...
NameMap name_map = get(vertex_name, G);
name_map[u] = "Hans";
```

- `name_map` is just a handle (pointer), copying it costs  $O(1)$ .
- `vertex_name_t` is a predefined tag. It is purely conventional (you can create `property<vertex_name_t, int>` and store distances), but algorithms use them as default choices if not instructed otherwise.

## Useful stuff: Edge property map

```
typedef adjacency_list<vecS, vecS, directedS,  
    no_property,    // Vertex properties, none this time.  
    // Edge properties as fifth template argument.  
    property<edge_weight_t, int> > > Graph;  
typedef property_map<Graph, edge_weight_t>::type  
    WeightMap;  
...  
EdgeDescriptor e;  
...  
WeightMap wm = get(edge_weight, G);  
wm[e] = k;
```

- To close nested templates > > must be used instead of >>.

## Useful stuff: Some predefined properties

- `vertex_name_t`
- `vertex_distance_t`
- `vertex_color_t`
- `vertex_degree_t`
- `edge_name_t`
- `edge_weight_t`

Do not be misled into, e.g., thinking that `vertex_degree_t` will automatically keep track of the degree for you.



## Useful stuff: Defining a custom property

Convenient e.g., if you want to keep additional info associated with edges.

```
namespace boost {  
    enum edge_info_t { edge_info = 219 }; // A unique ID.  
    BOOST_INSTALL_PROPERTY(edge, info);  
}  
struct EdgeInfo {  
    ...  
};
```

## Useful stuff: Defining a custom property (cont.)

```
typedef adjacency_list<vecS, vecS, directedS,  
    no_property,  
    property<edge_info_t, EdgeInfo> > Graph;  
typedef property_map<Graph, edge_info_t>::type InfoMap;  
...  
InfoMap im = get(edge_info, G);  
im[e] = ...
```

## Useful stuff: Calling an algorithm in different ways

- Example: `kruskal_minimum_spanning_tree`.
- Follow the [doc page](#).
- Header: `#include <boost/graph/kruskal_min_spanning_tree.hpp>`

```
vector<Edge> mst;           // edge list to store mst
kruskal_minimum_spanning_tree(G, back_inserter(mst));
vector<Edge>::iterator ebegin, eend = mst.end();
// Go through minimum spanning tree
for (ebegin = mst.begin(); ebegin != eend; ++ebegin) {
    ...
}
```

- `edge_weight_t` map must be defined for this to work.

## Useful stuff: Calling an algorithm in different ways – named parameters

- Maybe you might want to access additional information computed by the Union-Find algorithm of Kruskal. For example you can access `rank_map` and `predecessor_map`.
- You need to provide additional custom arguments. This is done via *named parameters*.

```
vector<int> R(num_vertices(G)), P(num_vertices(G));  
kruskal_minimum_spanning_tree(G, back_inserter(mst),  
    weight_map(get(edge_weight, G)).  
        rank_map(&R[0])). // A dot, not a comma!  
    predecessor_map(&P[0]));
```

- To pass a vector as a vertex property map you need to provide `&V[0]`, an iterator like `V.begin()` will not work.
- Defaults: `get(edge_weight, G)` for `weight_map`, internally created vectors for `rank_map` and `predecessor_map`.